




**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# Domain-specific Libraries for Weather and Climate


CSCS User Lab Day

Lukas Mosimann, CSCS

September 09, 2019

# Weather and Climate Applications



- Application properties
  - Numerical methods
    - Explicit methods in the horizontal
    - Implicit methods in the vertical
  - Usually many inputs for one output
  - Mostly memory-bound
  - Distributed memory



$$\frac{\partial u}{\partial t} = \alpha \frac{\partial^2 u}{\partial x^2} \xrightarrow{\text{explicitly discretized}} \boxed{\begin{array}{l} \text{time step} \\ u_i^{n+1} = u_i^n + \alpha \frac{\Delta t}{(\Delta x)^2} (u_{i+1}^n - 2u_i^n + u_{i-1}^n) \\ \text{field} \qquad \qquad \qquad \text{spatial step} \end{array}}$$


# Algorithmic Patterns

- Domain is a cartesian grid
  - 2D / 3D stencil-like** computations
  - Embarrassingly parallel on the horizontal plane ( $I$ - $J$ )
  - Dependencies on the vertical ( $K$ )
    - Forward, backward or parallel
    - Partial specializations for intervals
  - Halo-updates / boundary conditions



Horizontal stencil

```
for (int i = 0; i < ni; ++i)
  for (int j = 0; j < nj; ++j)
    for (int k = 0; k < nk; ++k)
      out[i][j][k] = -4.0 * in[i][j][k]
        + in[i-1][j][k] + in[i+1][j][k]
        + in[i][j-1][k] + in[i][j+1][k];
```



Sequential vertical solver

```
for (int i, j = ...)
  c[i][j][0] /= b[i][j][0]
  for (int k = 1; k < nk; ++k)
    c[i][j][k] /= b[i][j][k]
      - a[i][j][k] * c[i][j][k-1]
```

# Challenges

- Traditional weather and climate codes:
  - Large monolithic codebases hand-tuned for a specific architecture
  - Compute motifs hidden behind these optimizations
  - Hardware-specific acceleration usually added using pragmas, e.g. OpenMP or OpenACC
- Current hardware trends
  - Different architectures require very specific optimizations
- **Problem:** Lack of *domain* specific libraries to write models in a performance portable way for multiple architectures
- **Solution:** GridTools C++ Libraries



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# GridTools: Libraries for Applications on Grids

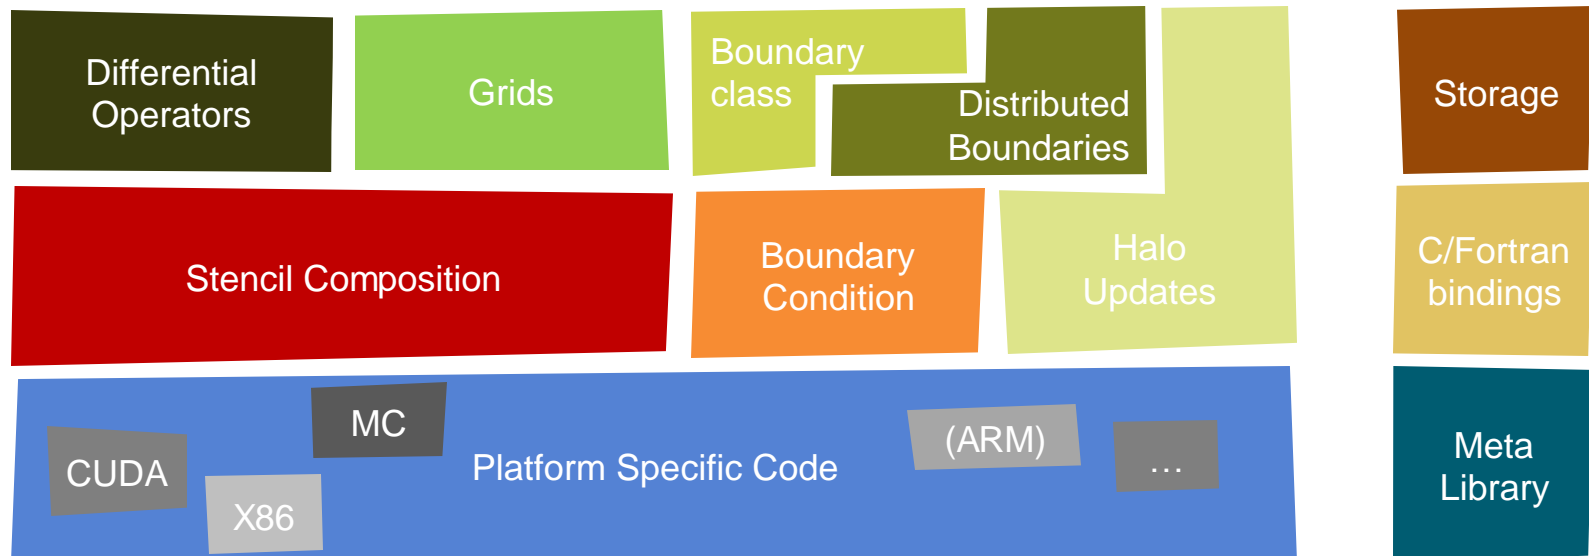

---

# GridTools: Libraries for Applications on Grids

- Open Source C++ project

- Goals


- Productivity, portability, performance
- Integration with existing codebases, e.g. COSMO (**Fortran**)



<https://www.github.com/GridTools/gridtools>

# Declarative Programming in GridTools

- Separation of concerns between interface and implementation
  - Abstract the underlying execution architecture / strategy
  - Provide enough information about **what** to do
  - Let the implementation choose **how** to do it
- Storages
  - Layout optimized for access patterns



# Stencil Composition in GridTools


- What is a stencil from a programming point of view?

- Loop control structures
- Loop body

```
for (int i = 0; i < ni; ++i)
  for (int j = 0; j < nj; ++j)
    for (int k = 0; k < nk; ++k)
      out[i][j][k] = -4.0 * in[i][j][k]
        + in[i-1][j][k] + in[i+1][j][k]
        + in[i][j-1][k] + in[i][j+1][k];
```

- How do we achieve good performance for the GridTools execution model?


- Naïve implementations are not efficient
- Hardware dependent iterations
- Loop fusion
  - Difficult to maintain manually
  - Hardware dependent





# Stencil Composition Example

- Stencil operator: Laplacian



- Composing stencils

```
struct laplacian {  
    using in = in_accessor<0, extent<-1, 1, -1, 1>>;  
    using lap = inout_accessor<1>;  
    using param_list = make_param_list<in, lap>;  
  
    template <typename Evaluation>  
    GT_FUNCTION static void apply(Evaluation const &eval) {  
        eval(lap(i, j, k)) = -4. * eval(in(i, j, k))  
            + eval(in(i + 1, j, k))  
            + eval(in(i, j + 1, k))  
            + eval(in(i - 1, j, k))  
            + eval(in(i, j - 1, k));  
    }  
};
```

```
using arg_in = arg<0, data_store_t>;  
using arg_lap = tmp_arg<1, data_store_t>;  
using arg_out = arg<2, data_store_t>;  
auto horizontal_diffusion = make_computation<backend_t>(  
    my_grid,  
    make_multistage(  
        execute::parallel(),  
        make_stage<laplacian>(arg_in(), arg_lap()),  
        make_stage<laplacian>(arg_lap(), arg_out())  
    ));  
horizontal_diffusion.run(arg_in{} = in, arg_out{} = out);
```





**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# GridTools4Py: Python Interface to GridTools

---

# GT4Py Goals

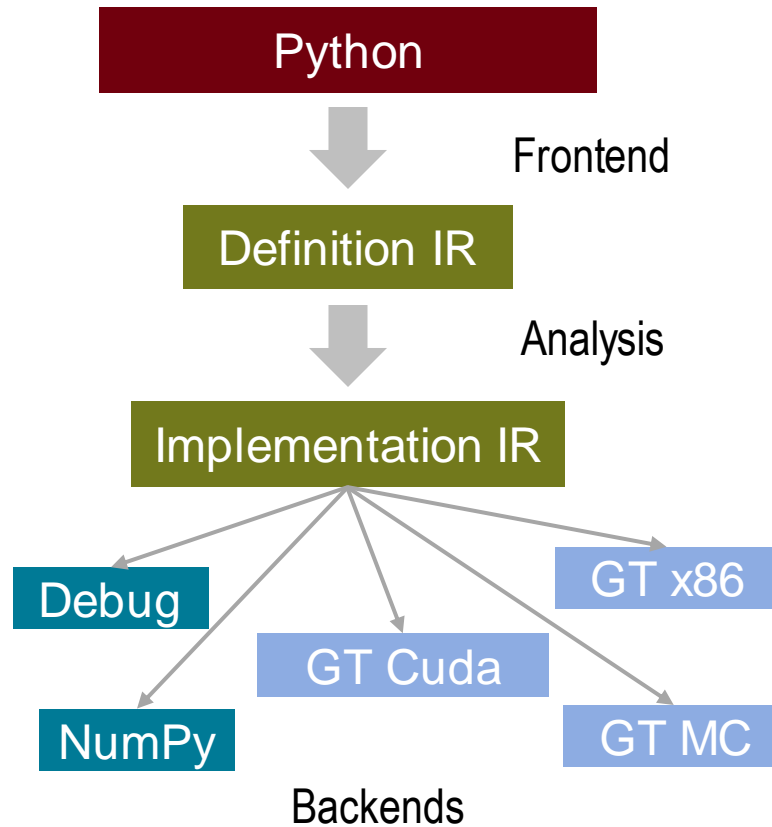
- C++ is less known in the community (Fortran, Python)
  - Too low-level for domain scientists
- **Solution:** Use Python as a higher level embedded DSL
  - Hide C++ templates boilerplate from the user
- Use GridTools benefits
  - Employ flexible execution model with multiple backends
  - Provide same performance / low overhead
- Raise level of abstraction
  - Automatic computation of
    - Halo extents and data dependencies
    - Stages and multi-stages / computation-on-the-fly
    - Temporary fields vs scalar variables
- First release in the coming weeks



# Horizontal Stencil Example


```
1 import gridtools as gt
2
3 backend = "gtx86"
4
5 # Storage definition
6 float_st = gt.storage.StorageDescriptor(dtype=np.float32, grid_group="main")
7
8 # Stencil definition
9 @gt.stencil(backend=backend)
10 def horizontal_diffusion(u: float_st, diffusion: float_st, coeff: float):
11     laplacian = 4.0 * u[0, 0, 0] - (u[1, 0, 0] + u[-1, 0, 0] + u[0, 1, 0] + u[0, -1, 0])
12     flux_i = laplacian[1, 0, 0] - laplacian[0, 0, 0]
13     flux_i = 0.0 if flux_i * (u[1, 0, 0] - u[0, 0, 0]) > 0 else flux_i
14     flux_j = laplacian[0, 1, 0] - laplacian[0, 0, 0]
15     flux_j = 0.0 if flux_j * (u[0, 1, 0] - u[0, 0, 0]) > 0 else flux_j
16     diffusion = u[0, 0, 0] - coeff * (
17         flux_i[0, 0, 0] - flux_i[-1, 0, 0] + flux_j[0, 0, 0] - flux_j[0, -1, 0]
18     )
19
20 # Storage initialization
21 # u_data_array = np.load(...)
22 u_storage = gt.storage.from_array(u_data_array,
23                                   descriptor=float_st,
24                                   backend=backend,
25                                   halo=(2, 2, 0))
26
27 out_storage = gt.storage.empty(u_data_array.shape,
28                                descriptor=float_st,
29                                backend=backend,
30                                halo=(2, 2, 0))
31
32 # Stencil execution
33 horizontal_diffusion(u_storage, out_storage, coeff=4.0)
34
35 print("Result: ", out_storage.data)
```

# Pipeline



# Preliminary Results

- Horizontal diffusion
  - GTMC and GTCuda backends (disabled caches)
  - Python overhead becomes neglectable for medium-sized domains






**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



**Thank you for your attention.**