




CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



A Practical Introduction to CSCS HPC Infrastructure

CSCS User Lab Day, Hochschule Luzern, September 9th 2019

Luca Marsella, CSCS



CSCS

|



ETH zürich

Outline of the Presentation

- User Policies
 - Data retention
 - Fair usage
- Resources
 - Filesystems
 - Computing
- Running Jobs
 - The Slurm scheduler
 - Best practices
- Troubleshooting
 - Documentation and list of FAQ
 - How to submit a support request



CSCS office building in Lugano

User Policies

General Policies

The code of conduct outlines proper practices:

- **Access to Source Codes:** you agree to make codes available for support
- **Scientific Advisory Committee:** committee members must not be contacted
- **Acknowledgements:** you must acknowledge the use of CSCS resources in all publications related to your production with reference to the “*project ID ###*”

User Regulations define basic guidelines:

- **Accounts are personal** and sharing them is forbidden
- **ETH Zurich Acceptable Use Policy for Telematics Resources (“BOT”)**

Access to CSCS resources **may be revoked to users violating the policies**

Data Retention Policies

Data backup for **active projects**:

- Data in **users** and **project** folders is backed up (past 90 days)
- Data recovery is also possible with daily snapshots (past 7 days)
- Data will be removed **3 months** after the expiration of the project

As soon as a **project expires**:

- Data backup is **disabled** immediately
- **No data recovery** after the final data removal

No backup for data in **scratch**:

- No recovery in case of **accidental data loss**
- No recovery of data deleted due to the cleaning policy

Fair Usage Policies

Slurm:

- The job scheduler is a shared resource among users submitting jobs
- **Do not submit large numbers of jobs** and commands at the same time
- We will be forced to kill jobs and limit new submissions

Login nodes:

- **Running applications on login nodes is not allowed**
- Submit your simulations with the Slurm scheduler on compute nodes
- **Heavy processes** running on login nodes will be **terminated**

Resources

How to access the systems

You should have already obtained an [account at CSCS](#)

The front end Ela is accessible via **ssh** as **ela.cscs.ch**:

```
$ ssh ela.cscs.ch
```

- It provides a minimal Linux environment
- You can **ssh** the computing systems from Ela
- You can start an [External Data Transfer](#) with GridFTP from/to CSCS

```
$ ssh daint.cscs.ch
```

Please note the following:

- **No programming environments** on the front end system
- User scratch space is **not accessible** from Ela

Filesystems

	/scratch (Piz Daint)	/scratch (Clusters)	/users	/project	/store
Type	Lustre	GPFS	GPFS	GPFS	GPFS
Quota	Soft quota 1 M files	None	10 GB/user 100K files	Maximum 50K files/TB	Maximum 50K files/TB
Expiration	30 days	30 days	Account closure	End of the project	End of the contract
Data Backup	None	None	90 days	90 days	90 days
Access Speed	Fast	Fast	Slow	Medium	Slow
Capacity	8.8 PB	1.4 PB	86 TB	4.7 PB	3.6 PB

Soft quota:

- Soft quota on **scratch** to prevent excessive loads on the Lustre filesystem
- Quota reached: **warning at submit time, no job submission** allowed

/scratch filesystem

Fast workspace for running jobs:

- Designed for **performance** rather than reliability
- **Cleaning policy**: files **older than 30 days deleted** daily
- **No backup**: transfer data after job completion

Performance of Piz Daint scratch (Lustre filesystem):

- **Soft quota on inodes** (files and folders) **to avoid** large numbers of small files
- Occupancy impacts performance:
 - **> 60%**: we will ask you to **remove unnecessary data immediately**
 - **> 80%**: we will free up disk space **manually removing data**

All CSCS systems provide a scratch personal folder:

- the variable **\$SCRATCH** points to the user space

/users and /project storage

Shared parallel filesystems based on the IBM GPFS software:

- Accessible from the login nodes using native GPFS client
- Storage space for datasets, shared code or configuration scripts
- Better performance with larger files (archive small files with **tar**)

Users are NOT supposed to run jobs here:

- The emphasis is on **reliability over performance**
- All directories are **backed up** with GPFS snapshots
- No cleaning policy until **3-months** after the end of the project

Environment variables pointing to personal folders:

- **\$HOME** points to **/users/\$USER**
- **\$PROJECT** points to **/project/<group_id>/\$USER**

Computing Resources

Computing time on Cray systems accounted in **node hours**:

- Resources assigned over **three-months windows**
 - **Quotas reset** on April 1st, July 1st, October 1st and January 1st
- Use thoroughly the quarterly compute budget within the time frame
- Unused resources in the three-months periods **cannot be recovered**

Check your budget in the **current allocation window**:

- Group usage - **sbuccheck**
 - reports group usage across the various systems
- Daily usage - **monthly_usage**
 - **monthly_usage --individual** usage per group member
- Overview of resources with the [Account and Resources Tool](#)
 - Check the details on the [dedicated page](#) of the User Portal

Piz Daint Specifications

Model	Cray XC50/XC40
XC50 Compute Nodes (Intel Haswell processor)	Intel® Xeon® E5-2690 v3 @ 2.60GHz (12 cores, 64GB RAM) and NVIDIA® Tesla® P100 16GB
XC40 Compute Nodes (Intel Broadwell processor)	Intel® Xeon® E5-2695 v4 @ 2.10GHz (18 cores, 64/128 GB RAM)
Login Nodes	Intel® Xeon® E5-2650 v3 @ 2.30GHz (10 cores, 256 GB RAM)
Interconnect Configuration	Aries routing and communications ASIC Dragonfly network topology
Scratch capacity	Piz Daint scratch filesystem: 8.8 PB

File Systems:

- The variable **\$SCRATCH** points at user space **/scratch/snx3000/\$USER**
- **/project** and **/store** mounted with **read-only** access on compute nodes

Setting the Programming Environment

You should prepare the environment before running jobs:

- CSCS systems use the modules framework
 - It manages applications and libraries path
 - Check currently loaded modules with **module list**
- **Modules** loaded at login
 - The default environment on Piz Daint is **PrgEnv-cray**
 - The default architecture is XC50 (Intel Haswell): **craype-haswell**
 - Browse the available modules with **module avail**
- **Adjust targets** according to your project (**sbuchek**)
 - **daint-gpu** targets the XC50 (Intel Haswell and P100 Tesla GPUS)
 - **daint-mc** targets the XC40 (Intel Broadwell multicore)
 - These modules update the **MODULEPATH**: use **module switch** to swap



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich

Running Jobs

The Slurm scheduler

- Slurm is the batch system/scheduler running on CSCS machines
- Permits users to run jobs with specific settings
- Job submission by calling `sbatch` with a job script

job.sh

```
#!/bin/bash -l
#SBATCH --nodes=10
#SBATCH --time=0:30:00
#SBATCH --partition=normal
#SBATCH --constraint=gpu
[...]

srun myprogram
```

```
$ sbatch job.sh
```


Slurm jobscripts

- Our web interface for generating job scripts covers most cases

Slurm Jobscript Generator

https://user.cscs.ch/access/running/jobscript_generator/

GETTING STARTED

Accounting

Running Jobs

Jobscript Generator

Fulen

Grand Tuvé

Piz Daint

Technical Report

Slurm Jobscript Generator

Computing system

Select the computing system on which you want to submit your job.

Daint MultiCore

Partition

Select the partition on which you want to submit your job.

normal

Executable

- For a comprehensive list of all options check:

```
$ man sbatch
```

Slurm queues on Piz Daint

- Corresponding Slurm option: --partition
- Permits users to run jobs on different queues

Name of the queue	Max time	Max nodes	Brief Description
debug	30 min	4	Quick turnaround for test jobs (one per user)
large	12 h	4400	Large scale work, by arrangement only
long	72 h	4	Maximum 5 long jobs in total (one per user)
low	6 h	2400(gpu)/512(mc)	(currently disabled)
normal	24 h	2400(gpu)/512(mc)	Standard queue for production work
prepost	30 min	1	High priority pre/post processing
xfer	24h	1	Data transfer queue
total	2 h		CSCS maintenance queue (restricted use)



https://user.cscs.ch/access/running/piz_daint/#slurm-batch-queues

Slurm queues (2)

- Watch your jobs in queues with

```
$ squeue -u ${USER}
```

```
daint103:~$ squeue -u simbergm
```

JOBID	USER	ACCOUNT	NAME	ST	REASON	START_TIME	TIME	TIME_LEFT	NODES	CPUS
11942503	simbergm	csstaff	hpx-3662-gcc-7	R	None	16:36:57	30:26	5:29:34	1	24
11945966	simbergm	csstaff	hpx-3712-gcc-7	R	None	16:44:24	22:59	5:37:01	1	72
11947200	simbergm	csstaff	hpx-3229-clang	PD	BeginTime	17:34:15	0:00	6:00:00	1	1
11947180	simbergm	csstaff	hpx-3684-gcc-7	PD	BeginTime	Tomorr 00:19	0:00	6:00:00	1	1

- Observe state of queues with

```
$ sinfo -o"%P %.5a %.10l %.6D %.6t"
```

```
daint103:~$ sinfo
```

PARTITION	AVAIL	JOB_SIZE	TIMELIMIT	CPUS	S:C:T	NODES	STATE	NODELIST
debug	up	1-4	30:00	72	2:18:2	2	allocated	nid00[448-449]
debug	up	1-4	30:00	24+	1+:12+	14	idle	nid0[0008-0011,0450-0451,3508-3511,4276-4279]
xfer	up	1	1-00:00:00	9	9:1:1	2	idle	nordend0[3-4]
uftp	up	1	1-00:00:00	0	0:0:0	0	n/a	
cscsci	up	1	1-00:00:00	24+	1+:12+	7	down\$	nid0[0125,0299,3541-3543,4579,5967]
cscsci	up	1	1-00:00:00	24+	1+:1+:	28	maint	nid0[0124,0126,1144-1147,1804-1807,3492-3495,3576-

Job Priority

- Job priority is based on partition, fairshare and waiting time

```
$ sprio -w
```

- Check the reason why a job is pending with

```
$ squeue -u ${USER}
```

- Check your budget with

```
$ sbucheck
```

Even if you still have lots of hours left, there may be other users/accounts with less usage and/or more hours allocated

- If the reason is “priority”, then you have to wait longer!

- Also, make sure there are no reservations in the system (maintenances, large runs, etc.)

```
$ scontrol show reservations
```

Job allocations

- CSCS has 3-month allocation periods
- If you want to fully utilize your allocated node hours, it's better to have a constant stream of jobs rather than packing all the jobs at the end of the allocation period

Good practices when submitting jobs

Specify accurate wall time

```
#!/bin/bash -l
#SBATCH --nodes=120
#SBATCH --time=0:30:00
#SBATCH --partition=normal
#SBATCH --constraint=gpu
[...]
```

Run jobs off \${SCRATCH}

```
#!/bin/bash -l
#SBATCH --nodes=120
#SBATCH --time=0:30:00
#SBATCH --partition=normal
#SBATCH --constraint=gpu
#SBATCH --mail-type=ALL
#SBATCH --mail-user=<your_email>

cd ${SCRATCH}
srun ${SCRATCH}/my_binary
```

For jobs with *many* tasks, use greasy

```
#!/bin/bash -l
#SBATCH --nodes=120
#SBATCH --time=0:30:00
#SBATCH --partition=normal
#SBATCH --constraint=gpu

module load daint-mc
module load GREASY/2.1-cscs-CrayGNU-18.08
greasy -f greasy_tasks.list
```

Make sure your sruns work! (or sleep a little bit in between)

```
#!/bin/bash -l
#SBATCH --nodes=120
#SBATCH --time=0:30:00
#SBATCH --partition=normal
#SBATCH --constraint=gpu
function p() {
  rt=$?;
  if [[ ${rt} -ne 0 ]]; then
    sleep 2
  fi
  return ${rt}
}
srun mytask ; p
srun mytask2 ; p
```

What not to do when submitting jobs

Jobs that submit other jobs/tasks in loops

```
#!/bin/bash
#SBATCH ...
while :
do
  srun sbatch my_job.sbatch
  sleep 1
done
```

Jobs with *thousands* of tasks

```
# sacct -j 123456789 |wc -l
25337
```

Jobs that run off `${HOME}`

```
#!/bin/bash -l
#SBATCH --nodes=120
#SBATCH --time=0:30:00
#SBATCH --partition=normal
#SBATCH --constraint=gpu

srun ~/my_binary ~/Large_input
```

Jobs with hundreds of tasks in parallel

```
#!/bin/bash -l
#SBATCH --nodes=3
#SBATCH --time=0:30:00
#SBATCH --partition=normal
#SBATCH --constraint=gpu

export CRAY_CUDA_MPS=1
cd $SLURM_SUBMIT_DIR

date
srun --nodes=1 --bcast=/tmp/${USER} --ntasks=1 --ntasks-per-node=1 --cpus-
per-task=12 tune_5x16x13_exe0 &
srun --nodes=1 --bcast=/tmp/${USER} --ntasks=1 --ntasks-per-node=1 --cpus-
per-task=12 tune_5x16x13_exe1 &
srun --nodes=1 --bcast=/tmp/${USER} --ntasks=1 --ntasks-per-node=1 --cpus-
per-task=12 tune_5x16x13_exe10 &
[...]
sleep 29m
```

What not to do on login nodes

queue without filtering

```
$ queue | grep ${USER}
$ queue -u ${USER}
```



watch overloads the scheduler

```
$ watch queue -u ${USER}
```



sacct + watch: even worse!

```
$ watch sacct
```



GNU make without number of tasks

```
$ make -j
$ make -j8
```



Avoid infinite loops

```
#!/bin/bash
while :
do
clear
queue | grep JOBID
queue | grep ${USER}
sleep 1
done
```



Other loops are even more evil!

```
#!/bin/bash
for i in ${var}; do
sbatch my_job.sbatch
done
```



Use e-mail notification instead of loops with Slurm commands

```
#!/bin/bash -l
#SBATCH --nodes=2
#SBATCH --mail-type=ALL
#SBATCH --mail-user=<your_email>
```



Summary: How to submit jobs at CSCS

- Move input data to \$SCRATCH

```
$ cd $SCRATCH  
$ cp -r ~/input .
```

- Use the jobscript generator and accurately specify runtime

```
$ sbatch job.sh
```

- Monitor (manually) your jobs with

```
$ squeue -u ${USER}
```

- Use Slurm e-mail notification for live-updates on job status instead of repeated Slurm commands
- Copy important output data back to /project or /home



Troubleshooting

What to do in case of trouble?

If you experience an issue on the systems:

- ☐ Search the content of the [User Portal](#) (top right field)
- ☐ Does your issue match any [Frequently Asked Question](#)?
- ☐ For advanced topics, user's guides may also help
 - **module help**
 - **man**
 - **CrayPubs**

Basic Documentation

- User Portal: <http://user.cscs.ch>
- Frequently Asked Questions: <https://user.cscs.ch/access/faq>
- More info with **module help**:
 - **module help cce**
- Manuals and User's Guides: command **man** on the shell

```
$ module help PrgEnv-cray
```

```
----- Module Specific Help for 'PrgEnv-cray/6.0.4' -----
```

```
The PrgEnv-cray modulefile loads the Cray Programming Environment, which includes the Cray Compiling Environment (CCE). This modulefile defines the system paths and environment variables needed to build an application using CCE for supported Cray systems. For more information on using targeting modules see Cray Programming Environment User's Guide, S-2529-114.
```

```
This module loads the following products:
```

```
craype  
cce  
cray-libsci  
udreg  
ugni  
pmi  
dmapp  
gni-headers  
xpmem  
job  
dvs  
alps  
rca  
atp  
perftools-base
```

Cray Documentation

- **CrayPubs** at <http://pubs.cray.com>:
 - Quick access and search of Cray books
 - **man** pages and third-party documentation
 - Available in HTML and PDF formats
- **Cray man** pages:
 - Textual help files on the command line of Cray systems
 - **man** command followed by the name of the man page
 - Described on man(1) page accessible with **man man**

How to submit a support request

Contact us **if you can't find a solution:**

- ❑ Write an e-mail to help@cscs.ch
- ❑ Specify the **system** and your **project ID** in the subject
- ❑ Report the **Slurm job ID** and indicate the **Slurm job script**
- ❑ Copy scripts and source files to **\$SCRATCH** and give us access

The **more detailed the request**, the **more effective the reply!**

Example of a request for support

Template message to be adapted for help@cscs.ch:

Subject: Slurm job failed on Piz Daint (project *<project ID>*)

Content:

My username is *<user name>*, I submitted the job *<job ID>* on Piz Daint.

The job running *<code name>* exited with state **FAILED** but no error in output.

The job script (*<script name>*) and input files (*<file list>*) can be found here:

- **/scratch/snx3000/\$USER/job**
- I have given read access with **chmod -R +r \$SCRATCH/job**

Please let me know the reason of the failure.

Useful links

- CSCS User Portal:
 - <https://user.cscs.ch>
- Cray Documentation:
 - <https://pubs.cray.com>
- NVIDIA Documentation:
 - <https://docs.nvidia.com>
- Contact us:
 - help@cscs.ch



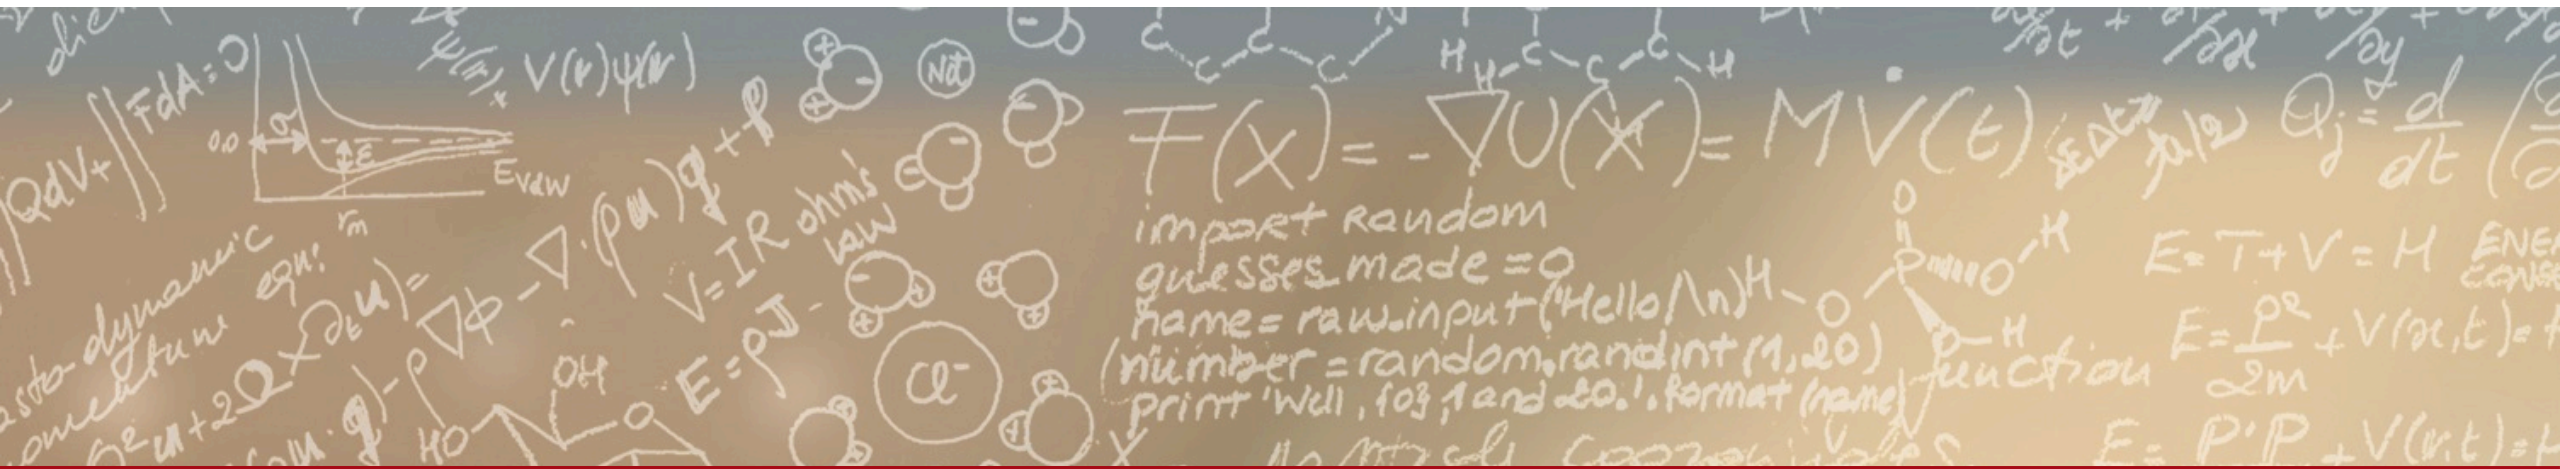
Piz Daint in the machine room at CSCS



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



Thank you for your kind attention