




**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# A Practical Introduction to CSCS HPC Infrastructure

CSCS User Lab Day, August 31<sup>st</sup> 2020

Luca Marsella, CSCS

# Outline of the Presentation

- Policies and Resources
  - Login
  - Filesystems
- Development Environment
  - Features of the hybrid system
  - Programming Environment
- Job Submission and Monitoring
  - Slurm workload manager
  - Best practices
- Documentation and Troubleshooting
  - How to submit a support request



*CSCS office building in Lugano*



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# Policies and Resources

---

# General Policies

The code of conduct outlines proper practices:

- **Access to Source Codes:** you agree to make codes available for support
- **Scientific Advisory Committee:** committee members must not be contacted
- **Acknowledgements:** you must acknowledge the use of CSCS resources in all publications related to your production with reference to the “*project ID ###*”

User Regulations define basic guidelines:

- **Accounts are personal** and sharing them is forbidden
- **Data ownership:** access to and use of data of other accounts without prior consent from the principal investigator is strictly prohibited
- **ETH Zurich Acceptable Use Policy for Telematics Resources (“BOT”)**

Access to CSCS resources **may be revoked to users violating the policies**

# Data Retention Policies

## Data backup for **active projects**:

- Data in **users** and **project** folders is backed up (past 90 days)
- Data recovery is also possible with daily snapshots (past 7 days)
- Data in **project** folders removed **3 months** after the **end of the project**

## As soon as a the project **expires**:

- Data backup is **disabled** immediately
- **No data recovery** after the final data removal

## **No backup** for data in **scratch**:

- No recovery in case of **accidental data loss**
- No recovery of data deleted due to the cleaning policy

# Fair Usage Policies

## Slurm:

- The job scheduler is a shared resource among users submitting jobs
- **Do not submit large numbers of jobs** and commands at the same time
- We will be forced to kill jobs and limit new submissions

## Login nodes:

- **Running applications on login nodes is not allowed**
- Submit your simulations with the Slurm scheduler on compute nodes
- **Heavy processes** running on login nodes will be **terminated**

Please check the summary of CSCS policies at  
<https://user.cscs.ch/access/accounting/#policies>

# How to access the systems

You should have already obtained an [account at CSCS](#)

The front end Ela is accessible via **ssh** as **ela.cscs.ch**:

```
$ ssh ela.cscs.ch
```

- It provides a minimal Linux environment
- You can **ssh** the computing systems from Ela
- You can start an [External Data Transfer](#) with GridFTP from/to CSCS

```
$ ssh daint.cscs.ch
```

Please note the following:

- **No programming environments** on the front end system
- User scratch space is **not directly accessible** from Ela

# Filesystems

	/scratch (Piz Daint)	/scratch (Clusters)	/users	/project	/store
Type	Lustre	GPFS	GPFS	GPFS	GPFS
Quota	Soft quota 1 M files	None	50 GB/user 500K files	Maximum 50K files/TB	Maximum 50K files/TB
Expiration	30 days	30 days	Account closure	End of the project	End of the contract
Data Backup	None	None	90 days	90 days	90 days
Access Speed	Fast	Fast	Slow	Medium	Slow
Capacity	8.8 PB	1.9 PB	160 TB	6.3 PB	5.0 PB

## Soft quota:

- Soft quota on **Piz Daint** to prevent excessive loads on the scratch filesystem
- Quota reached: **warning at submit time, no job submission** allowed



# /scratch filesystem

Fast workspace for running jobs:

- Designed for **performance** rather than reliability
- **Cleaning policy**: files **older than 30 days deleted** daily
- **No backup**: transfer data after job completion

Performance of Piz Daint scratch (Lustre filesystem):

- **Soft quota on inodes** (files and folders) **to avoid** large numbers of small files
- Occupancy impacts performance:
  - **> 60%**: we will ask you to **remove unnecessary data immediately**
  - **> 80%**: we will free up disk space **manually removing data**

All CSCS systems provide a scratch personal folder:

- the variable **\$SCRATCH on Piz Daint** points to **/scratch/snx3000/\$USER**

## /users and /project filesystems

Shared parallel filesystems based on the IBM GPFS software:

- Accessible from the login nodes using native GPFS client
- Storage space for datasets, shared code or configuration scripts
- Better performance with larger files (archive small files with **tar**)

**Users are NOT supposed to run jobs here:**

- The emphasis is on **reliability over performance**
- All directories are **backed up** with GPFS snapshots
- No cleaning policy until 3-months after the end of the project

Environment variables pointing to personal folders:

- **\$HOME** points to **/users/\$USER**
- **\$PROJECT** points to **/project/<group\_id>/\$USER**

# Computing Resources

Computing time on Cray systems is accounted in **node hours**:

- Resources are assigned over **three-months windows**
  - **Quotas reset** on April 1st, July 1st, October 1st and January 1st
- Use thoroughly the quarterly compute budget within the time frame
- Unused resources in the allocation periods **cannot be recovered**

Check your budget in the **current allocation window**:

- Group usage - **sbuccheck**
  - reports group usage across the various systems
- Daily usage - **monthly\_usage**
  - **monthly\_usage --individual** usage per group member
- Overview of resources with the [Account and Resources Tool](#)
  - Check the details on the [dedicated page](#) of the User Portal

# Development Environment

---

# Piz Daint Specifications

Model	Cray XC50/XC40
XC50 Compute Nodes (Intel Haswell processor)	Intel® Xeon® E5-2690 v3 @ 2.60GHz (12 cores, 64GB RAM) and NVIDIA® Tesla® P100 16GB
XC40 Compute Nodes (Intel Broadwell processor)	Intel® Xeon® E5-2695 v4 @ 2.10GHz (2 x 18 cores, 64/128 GB RAM)
Login Nodes	Intel® Xeon® E5-2650 v3 @ 2.30GHz (10 cores, 256 GB RAM)
Interconnect Configuration	Aries routing and communications ASIC Dragonfly network topology
Scratch capacity	Piz Daint scratch filesystem: 8.8 PB

## File Systems:

- **/project** is mounted with **read-only** access on compute nodes

# NVIDIA CUDA Toolkit v10.1

- Comprehensive development environment to build **GPU-accelerated applications**
  - **compiler** for NVIDIA GPUs
  - optimized **math libraries**
  - **debugging** and **performance** tools
- Features **programming guides**, user manuals, API reference and **online documentation** to get started quickly
- NVIDIA developer portal:  
<https://developer.nvidia.com/cuda-zone>



*NVIDIA Tesla P100 GPU Accelerator*

# Cray Linux Environment 7.0 UP01

- Cray Linux Environment (CLE) is the operating system on Cray systems
- CLE 7.0 UP01 is based on the SUSE Linux Enterprise Server version 15
- CLE 7.0 UP01 software release is available on Piz Daint since Nov 2019
- Read more on the [Cray Pubs Portal](#)



# Setting the Programming Environment

You should prepare the environment before running jobs:

- CSCS systems use the modules framework
  - The modules manage applications and libraries path
  - You can check currently loaded modules with **module list**
- Some **modules** are already loaded at login
  - The default environment on Piz Daint is **PrgEnv-cray**
  - The default architecture is XC50 (Intel Haswell): **craype-haswell**
  - You can browse the available modules with **module avail**
- You must **adjust your target architecture** (see **sbucheck**)
  - **daint-gpu** targets the XC50 (Intel Haswell and P100 Tesla GPUS)
  - **daint-mc** targets the XC40 (Intel Broadwell multicore)
  - These modules update the **MODULEPATH**



# Setting the Programming Environment

```
$ module switch PrgEnv-cray/6.0.5 PrgEnv-gnu
```

```
$ module load daint-gpu
```

```
$ module list
```

Currently Loaded Modulefiles:

1) modules/3.2.11.3	9) cray-libsci/19.06.1	17) dvs/2.12_2.2.151-7.0.1.1_5.6__g7eb5e703
2) gcc/8.3.0	10) udreg/2.3.2-7.0.1.1_3.9__g8175d3d.ari	18) alps/6.6.56-7.0.1.1_4.10__g2e60a7e4.ari
3) craype-haswell	11) ugni/6.0.14.0-7.0.1.1_7.10__ge78e5b0.ari	19) rca/2.2.20-7.0.1.1_4.9__g8e3fb5b.ari
4) craype-network-aries	12) pmi/5.0.14	20) atp/2.1.3
5) craype/2.6.1	13) dmapp/7.1.1-7.0.1.1_4.8__g38cf134.ari	21) perftools-base/7.1.1
6) cray-mpich/7.7.10	14) gni-headers/5.0.12.0-7.0.1.1_6.7__g3b1768f.ari	22) PrgEnv-gnu/6.0.5
7) slurm/19.05.3-2	15) xpmem/2.2.19-7.0.1.1_3.7__gdcf436c.ari	23) daint-gpu
8) xalt/2.7.24	16) job/2.2.4-7.0.1.1_3.8__g36b56f4.ari	

```
$ module avail ...
```

# Cray XC Programming Environment

- Cray XC PE 19.10 includes the **Cray Developer Toolkit - CDT 19.10**
  - non-default Programming Environments can be accessed with **cdt** modules
- The following products have been updated within this release:
  - **Cray Compiling Environment - CCE**
    - cce 9.0.2, cray-mpich 7.7.10, cray-libsci 19.06.1
  - **Cray Performance Measurement & Analysis Tools - CPMAT**
    - Perftools 7.1.1
  - **Cray Environment Setup and Compiling support - CENV**
    - cray-modules 3.2.11.3 and craype 2.6.1
  - **Third party products**
    - GCC 7.3.0 and 8.3.0, cray-python 2.7.15.7 and 3.6.5.7, cray-R 3.4.2

# Main default modules for supported applications on Piz Daint

Amber/18-14-14-CrayGNU-19.10-cuda-10.1

Boost/1.70.0-CrayGNU-19.10-python3

CDO/1.9.5-CrayGNU-19.10

CP2K/6.1-CrayGNU-19.10-cuda-10.1

CPMD/4.1-CrayIntel-19.10

GROMACS/2018.6-CrayGNU-19.10-cuda-10.1

LAMMPS/22Aug18-CrayGNU-19.10-cuda-10.1

NAMD/2.13-CrayIntel-19.10-cuda-10.1

NCL/6.4.0

NCO/4.8.1-CrayGNU-19.10

ParaView/5.7.0-CrayGNU-19.10-EGL

QuantumESPRESSO/6.4.1-CrayIntel-19.10-cuda-10.1

VASP/5.4.4-CrayIntel-19.10-cuda-10.1

Amber/18-14-14-CrayGNU-19.10

Boost/1.70.0-CrayGNU-19.10-python3

CDO/1.9.5-CrayGNU-19.10

CP2K/6.1-CrayGNU-19.10

CPMD/4.1-CrayIntel-19.10

GROMACS/2018.6-CrayGNU-19.10

LAMMPS/22Aug18-CrayGNU-19.10

NAMD/2.13-CrayIntel-19.10

NCL/6.4.0

NCO/4.8.1-CrayGNU-19.10

ParaView/5.7.0-CrayGNU-19.10-OSMesa

QuantumESPRESSO/6.4.1-CrayIntel-19.10

VASP/5.4.4-CrayIntel-19.10

Visit/3.1.0-CrayGNU-19.10

For more details please check the User Portal at <https://user.cscs.ch/computing/applications>



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# Job Submission and Monitoring

---

# The Slurm workload manager

- Slurm is the batch system/scheduler running on CSCS machines
- Permits users to run jobs with specific settings
- Job submission by calling sbatch with a job script

job.sh

```
#!/bin/bash -l  
#SBATCH --nodes=10  
#SBATCH --time=0:30:00  
#SBATCH --partition=normal  
#SBATCH --constraint=gpu  
[...]
```

```
srun myprogram
```

```
$ sbatch job.sh
```

# Slurm jobscripts

- Our web interface for generating job scripts covers most cases

Slurm Jobscript Generator

[https://user.cscs.ch/access/running/jobscript\\_generator/](https://user.cscs.ch/access/running/jobscript_generator/)

GETTING STARTED

Accounting

Running Jobs

Jobscript Generator

Fulen

Grand Tuvé

Piz Daint

Technical Report

Slurm Jobscript Generator

**Computing system**

Select the computing system on which you want to submit your job.

Daint MultiCore

**Partition**

Select the partition on which you want to submit your job.

normal

**Executable**

- For a comprehensive list of all options check:

```
$ man sbatch
```

# Slurm queues on Piz Daint

- Corresponding Slurm option: --partition
- Permits users to run jobs on different queues

Name of the queue	Max time	Max nodes	Brief Description
debug	30 min	4	Quick turnaround for test jobs (one per user)
large	12 h	4400	Large scale work, by arrangement only
long	7 days	4	Maximum 5 long jobs in total (one per user)
low	6 h	2400(gpu)/512(mc)	Currently disabled
<b>normal</b>	<b>24 h</b>	<b>2400(gpu)/512(mc)</b>	<b>Standard queue for production work</b>
prepost	30 min	1	High priority pre/post processing
xfer	24h	1	Data transfer queue

More information at [https://user.cscs.ch/access/running/piz\\_daint/#slurm-batch-queues](https://user.cscs.ch/access/running/piz_daint/#slurm-batch-queues)

## Slurm queues (2)

- Watch your jobs in queues with

```
$ squeue -u ${USER}
```

```
daint103:~$ squeue -u simbergm
```

JOBID	USER	ACCOUNT	NAME	ST	REASON	START_TIME	TIME	TIME_LEFT	NODES	CPUS
11942503	simbergm	csstaff	hpx-3662-gcc-7	R	None	16:36:57	30:26	5:29:34	1	24
11945966	simbergm	csstaff	hpx-3712-gcc-7	R	None	16:44:24	22:59	5:37:01	1	72
11947200	simbergm	csstaff	hpx-3229-clang	PD	BeginTime	17:34:15	0:00	6:00:00	1	1
11947180	simbergm	csstaff	hpx-3684-gcc-7	PD	BeginTime	Tomorr 00:19	0:00	6:00:00	1	1

- Observe state of queues with

```
$ sinfo -o"%P %.5a %.10l %.6D %.6t"
```

```
daint103:~$ sinfo
```

PARTITION	AVAIL	JOB_SIZE	TIMELIMIT	CPUS	S:C:T	NODES	STATE	NODELIST
debug	up	1-4	30:00	72	2:18:2	2	allocated	nid00[448-449]
debug	up	1-4	30:00	24+	1+:12+	14	idle	nid0[0008-0011,0450-0451,3508-3511,4276-4279]
xfer	up	1	1-00:00:00	9	9:1:1	2	idle	nordend0[3-4]
uftp	up	1	1-00:00:00	0	0:0:0	0	n/a	
cscsci	up	1	1-00:00:00	24+	1+:12+	7	down\$	nid0[0125,0299,3541-3543,4579,5967]
cscsci	up	1	1-00:00:00	24+	1+:1+	28	maint	nid0[0124,0126,1144-1147,1804-1807,3492-3495,3576-



# Job Priority

- Job priority is based on partition, fair share and waiting time `$ sprio -w`
- Check the reason why a job is pending with `$ squeue -u ${USER}`
- Check your budget with `$ sbuccheck`  
Even if you still have lots of hours left, there may be other users/accounts with less usage and/or more hours allocated
- If the reason is “priority”, then you have to wait in the queue!
- Also, make sure there are no reservations in the system (maintenances, large runs, etc.)

```
$ scontrol show reservations
```

# Good practices when submitting jobs

## Specify accurate wall time

```
#!/bin/bash -l
#SBATCH --nodes=120
#SBATCH --time=0:30:00
#SBATCH --partition=normal
#SBATCH --constraint=gpu
[...]
```

## Run jobs off \${SCRATCH}

```
#!/bin/bash -l
#SBATCH --nodes=120
#SBATCH --time=0:30:00
#SBATCH --partition=normal
#SBATCH --constraint=gpu
#SBATCH --mail-type=ALL
#SBATCH --mail-user=<your_email>

cd ${SCRATCH}
srun ${SCRATCH}/my_binary
```

## For jobs with *many* steps, use greasy

```
#!/bin/bash -l
#SBATCH --nodes=120
#SBATCH --time=0:30:00
#SBATCH --partition=normal
#SBATCH --constraint=gpu

module load daint-mc
module load GREASY
greasy -f greasy_tasks.list
```

## Make sure your sruns work! (or sleep a little bit in between)

```
#!/bin/bash -l
#SBATCH --nodes=120
#SBATCH --time=0:30:00
#SBATCH --partition=normal
#SBATCH --constraint=gpu
function p() {
  rt=$?;
  if [[ ${rt} -ne 0 ]]; then
    sleep 2
  fi
  return ${rt}
}
srun mytask ; p
srun mytask2 ; p
```

# What not to do when submitting jobs

## Jobs that submit other jobs/steps in loops

```
#!/bin/bash
#SBATCH ...
while :
do
srun sbatch my_job.sbatch
sleep 1
done
```

## Jobs with *thousands* of steps

```
# sacct -j 123456789 |wc -l
25337
```

## Jobs that run off \${HOME}

```
#!/bin/bash -l
#SBATCH --nodes=120
#SBATCH --time=0:30:00
#SBATCH --partition=normal
#SBATCH --constraint=gpu

srun ~/my_binary ~/Large_input
```

## Jobs with hundreds of steps in parallel

```
#!/bin/bash -l
#SBATCH --nodes=3
#SBATCH --time=0:30:00
#SBATCH --partition=normal
#SBATCH --constraint=gpu

export CRAY_CUDA_MPS=1
cd $SLURM_SUBMIT_DIR

date
srun --nodes=1 --bcast=/tmp/${USER} --ntasks=1 --ntasks-per-node=1 --cpus-per-task=12
tune_5x16x13_exe0 &
srun --nodes=1 --bcast=/tmp/${USER} --ntasks=1 --ntasks-per-node=1 --cpus-per-task=12
tune_5x16x13_exe1 &
srun --nodes=1 --bcast=/tmp/${USER} --ntasks=1 --ntasks-per-node=1 --cpus-per-task=12
tune_5x16x13_exe10 &
[...]
sleep 29m
```

# What not to do on login nodes

queue without filtering

```
$ queue | grep ${USER}
$ queue -u ${USER}
```



watch overloads the scheduler

```
$ watch queue -u ${USER}
```



sacct + watch: even worse!

```
$ watch sacct
```



GNU make without number of tasks

```
$ make -j
$ make -j8
```



Avoid infinite loops

```
#!/bin/bash
while :
do
clear
queue | grep JOBID
queue | grep ${USER}
sleep 1
done
```



Other loops are even more evil!

```
#!/bin/bash
for i in ${var}; do
sbatch my_job.sbatch
done
```



Use e-mail notification instead of loops with Slurm commands

```
#!/bin/bash -l
#SBATCH --nodes=2
#SBATCH --mail-type=ALL
#SBATCH --mail-user=<your_email>
```



# Summary: How to submit jobs at CSCS

- Move input data to \$SCRATCH

```
$ cd $SCRATCH  
$ cp -r ~/input .
```

- Use the jobscript generator and accurately specify runtime

```
$ sbatch job.sh
```

- Monitor (manually) your jobs with

```
$ squeue -u ${USER}
```

- Use Slurm e-mail notification for live-updates on job status instead of repeated Slurm commands
- Copy important output data back to /project or /home





**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

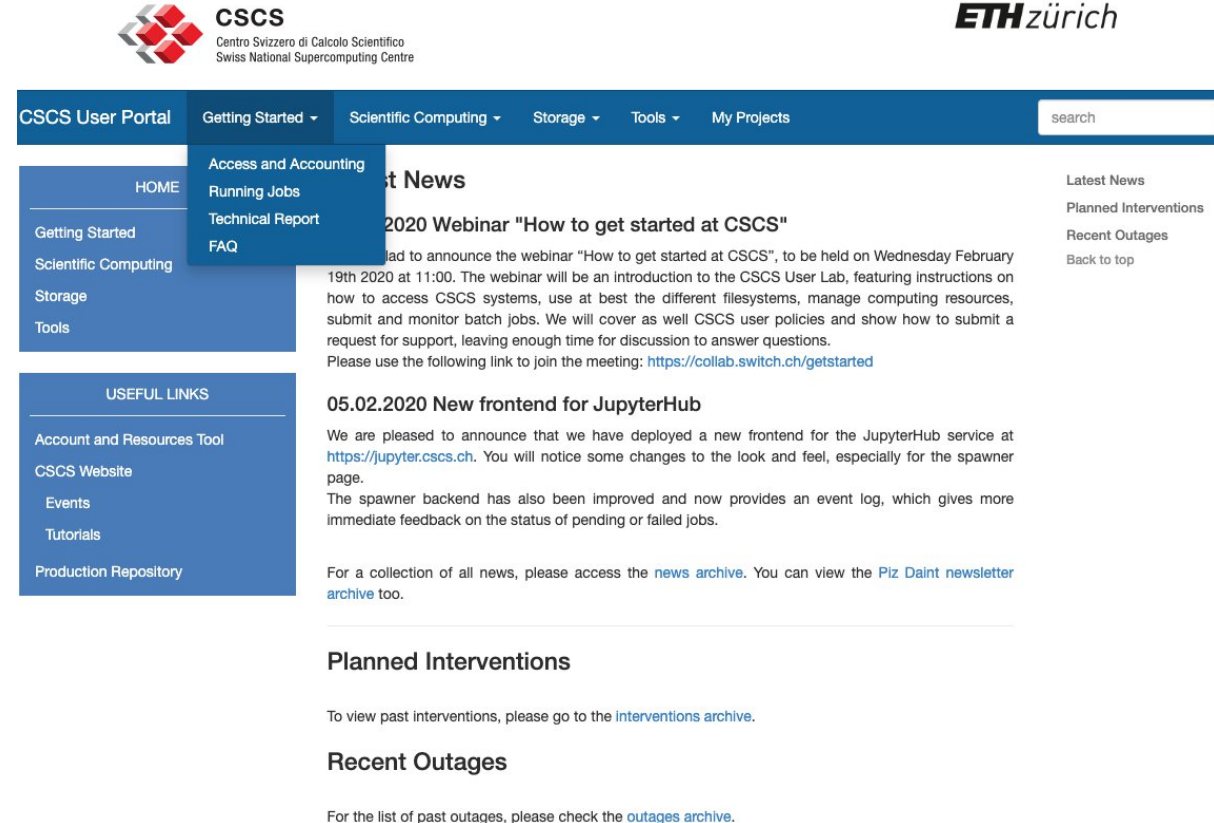
**ETH** zürich

# Documentation and Troubleshooting

---

# What to do in case of trouble?

- ❑ Search the content of the User Portal
- ❑ Check the Frequently Asked Questions
- ❑ Additional user documentation:
  - **module help**
  - **man**
  - **CrayPubs**



The screenshot shows the CSCS User Portal website. The header includes the CSCS logo (Centro Svizzero di Calcolo Scientifico / Swiss National Supercomputing Centre) and the ETH zürich logo. The navigation bar has links for Getting Started, Scientific Computing, Storage, Tools, and My Projects. A search bar is located on the right. The main content area features a sidebar with links to HOME, Getting Started, Scientific Computing, Storage, and Tools. The main content area displays the latest news, including a 2020 Webinar "How to get started at CSCS" and a new frontend for JupyterHub. The footer includes links to Account and Resources Tool, CSCS Website, Events, Tutorials, and Production Repository.

CSCS User Portal

Getting Started | Scientific Computing | Storage | Tools | My Projects

search

HOME

- Getting Started
- Scientific Computing
- Storage
- Tools

USEFUL LINKS

- Account and Resources Tool
- CSCS Website
- Events
- Tutorials
- Production Repository

Latest News

Planned Interventions

Recent Outages

Back to top

2020 Webinar "How to get started at CSCS"

ad to announce the webinar "How to get started at CSCS", to be held on Wednesday February 19th 2020 at 11:00. The webinar will be an introduction to the CSCS User Lab, featuring instructions on how to access CSCS systems, use at best the different filesystems, manage computing resources, submit and monitor batch jobs. We will cover as well CSCS user policies and show how to submit a request for support, leaving enough time for discussion to answer questions. Please use the following link to join the meeting: <https://collab.switch.ch/getstarted>

05.02.2020 New frontend for JupyterHub

We are pleased to announce that we have deployed a new frontend for the JupyterHub service at <https://jupyter.cscs.ch>. You will notice some changes to the look and feel, especially for the spawner page. The spawner backend has also been improved and now provides an event log, which gives more immediate feedback on the status of pending or failed jobs.

For a collection of all news, please access the [news archive](#). You can view the [Piz Daint newsletter archive](#) too.

Planned Interventions

To view past interventions, please go to the [interventions archive](#).

Recent Outages

For the list of past outages, please check the [outages archive](#).

# User Documentation

- General info with **module help**
  - **module help PrgEnv-cray**
- Product related details
  - **module help cce**
- Command specific guidelines
  - **man craycc**
- Advanced topics on CrayPubs
  - [CCE 9.1 Usage](#)

The PrgEnv-cray modulefile loads the Cray Programming Environment, which includes the Cray Compiling Environment (CCE). This modulefile defines the system paths and environment variables needed to build an application using CCE for supported Cray systems. For more information on using targeting modules see Cray Programming Environment User's Guide, S-2529-114.

This module loads the following products:

```
craype
cce
cray-libsci
udreg
ugni
pmi
dmapp
gni-headers
xpmem
job
dvs
alps
rca
atp
perftools-base
```



# Cray Documentation

- **Cray man pages:**
  - Textual help files on the command line of Cray systems
  - **man** command followed by the name of the man page
  - Described on man(1) page accessible with **man man**
- **CrayPubs** at <http://pubs.cray.com>:
  - Quick access and search of Cray books
  - **man** pages and third-party documentation
  - Available in HTML and PDF formats

# NVIDIA Documentation

- NVIDIA Documentation Portal
  - <http://docs.nvidia.com>
- Documentation on the system:
  - **module help cudatoolkit**
  - NVIDIA compiler
    - **nvcc --help**
  - CUDA profiler
    - **nvprof --help**

```
$ module help cudatoolkit

----- Module Specific Help for 'cudatoolkit/10.1.105_3.27-7.0.1.1_4.1__ga311ce7' -----

The modulefile defines the system paths and
variables for the Cuda Toolkit.

cray-cudatoolkit 10.1.105_3.27-7.0.1.1_4.1__ga311ce7
=====

Release Date: December 6, 2012

Purpose:
-----
cray-cudatoolkit 10.1.105_3.27-7.0.1.1_4.1__ga311ce7 provides a development environment
for NVIDIA GPUs. It provides a standalone compilation
environment for CUDA for C/C++ (nvcc) and the infrastructure
used by other compilation environments (CCE and PGI).

Includes:
  NVCC compiler
  CUDA runtime support libraries
  Ptx assembler
  CUDA math libraries
  Profiler
  Cuda-gdb
  Cuda-memcheck
```

# How to submit a support request

Contact us **if you can't find a solution:**

- ❑ Write an e-mail to [help@cscs.ch](mailto:help@cscs.ch)
- ❑ Specify the **system** and your **project ID** in the subject
- ❑ Report the **Slurm job ID** and indicate the **Slurm job script**
- ❑ Copy scripts and source files to **\$SCRATCH** and give us access

The **more detailed the request**, the **more effective the reply!**

# Example of a request for support

**Template** message to be adapted for [help@cscs.ch](mailto:help@cscs.ch)

**Subject:** Slurm job failed on Piz Daint (project *<project ID>*)

**Content:**

My username is *<user name>*, I submitted the job *<job ID>* on Piz Daint.

The job running *<code name>* exited with state **FAILED** but no error in output.

The job script (*<script name>*) and input files (*<file list>*) can be found here:

- **/scratch/snx3000/\$USER/job**
- I have given read access with **chmod -R +r \$SCRATCH/job**

Please let me know the reason of the failure.

# Useful links

- CSCS User Portal
  - <https://user.cscs.ch>
- Cray Documentation
  - <https://pubs.cray.com>
- NVIDIA Documentation
  - <https://docs.nvidia.com>
- Contact us
  - [help@cscs.ch](mailto:help@cscs.ch)



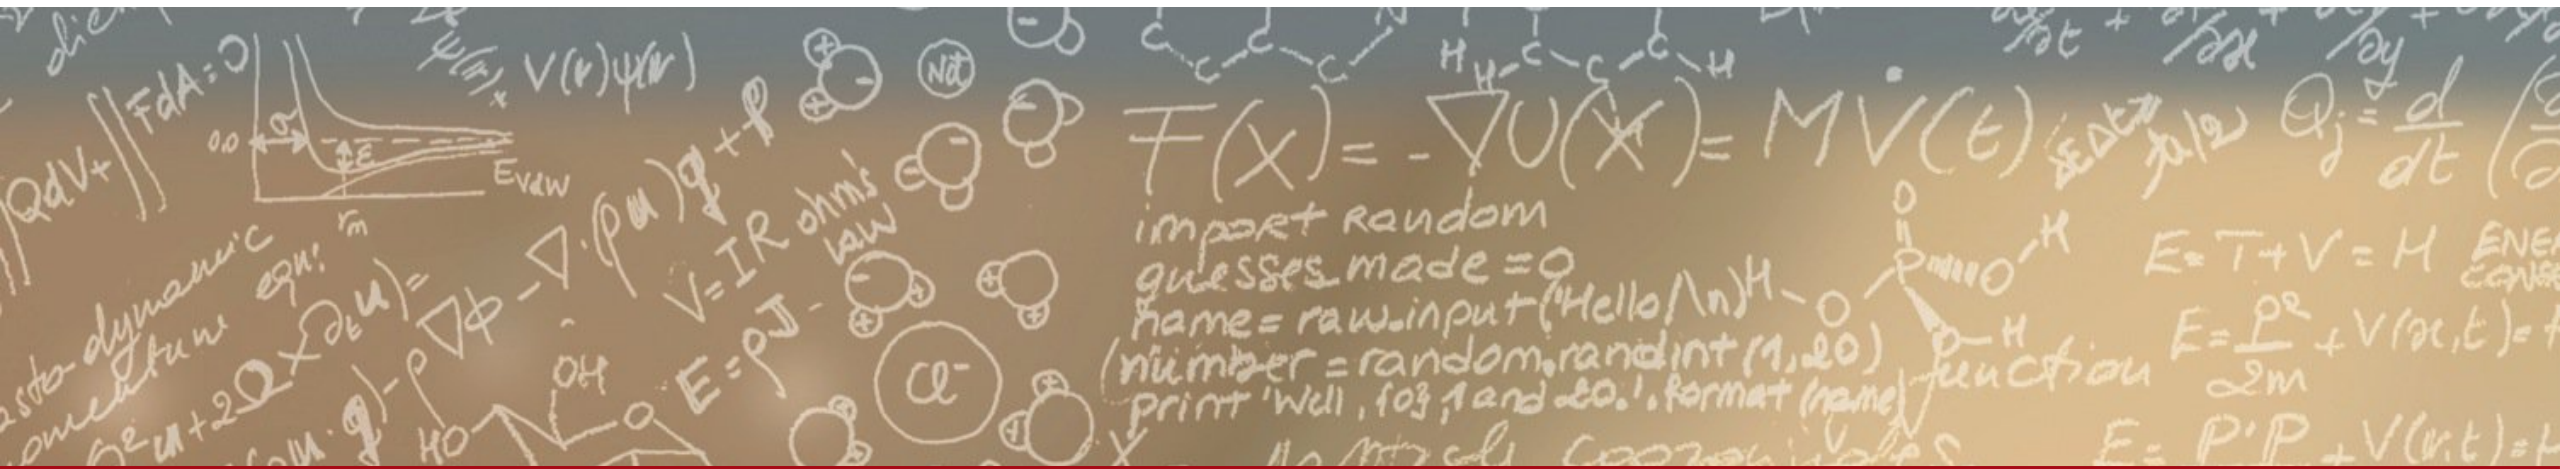
*Piz Daint in the machine room at CSCS*



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



**Thank you for your kind attention**